David J. Kimanthi

jephdavid@gmail.com linkedin.com/in/kimanthi
github.com/d-kimanthi United States

Professional Summary

Data Engineer with 8+ years of experience in software and data engineering. I have hands-on experience building scalable cloud data platforms and analytics systems on AWS, GCP, and Plantir Foundry. Expertise in designing reliable ELT pipelines, data warehouses, and data models using Python, Spark, SQL, dbt, and Airflow to support large-scale analytics and machine learning. Proven track record improving pipeline performance and delivering trusted datasets that enable product analytics, experimentation, and data-driven decision-making.

Skills

- **Programming:** Python, SQL, Java, FastAPI
- **Data Engineering & Processing:** PySpark, Apache Kafka, ETL/ELT, DBT, Batch & Streaming Data Processing
- **Data Warehousing & Modeling:** BigQuery, Snowflake, Databricks, DBT Core, Dimensional Modeling
- **Orchestration:** Apache Airflow, Dagster, Airbyte
- **Cloud/Data Platforms:** AWS, GCP , Palantir Foundry
- **Reporting & BI:** Power BI, Looker Studio, dbt Semantic Layer, KPI Dashboards
- **Data Quality & Governance:** Great Expectations, dbt Tests, pytest, Data Catalogs
- **Infrastructure & DevOps:** Docker, Terraform, Git, CI/CD Pipelines, Agile Development

Professional Experience

Proxify | Senior Data Engineer *Oct 2025 – Present*

- Architected and shipped a production-grade ELT pipeline ingesting **30K+ order line items daily** across 12 Amazon marketplaces using **SP-API, AWS S3, Snowflake,dbt Core**, powering executive revenue dashboards and Finance reporting.
- Designed a config-driven **Apache Airflow** framework with dynamic task mapping, region-scoped API rate limiting supporting both incremental and historical data backfill.
- Developed a **dbt Core** star schema (staging,intermediate, marts) implementing date based revenue recognition and dimensional models for sales, refunds, and settlement reconciliation.
- Partnered with Finance and business stakeholders to **define reporting requirements**, validate revenue business rules, and iteratively deliver trusted analytics solutions supporting executive decision-making.

IDinsight | Data Engineer III *Sep 2022 – Aug 2025*

- Architected and delivered an enterprise-grade data warehouse on AWS RDS (PostgreSQL) supporting **82+ KPIs across 8 national programs**, enabling centralized analytical reporting, scorecards, and data-driven decision-making for cross-functional teams.
- Eliminated **15+ hours** of weekly manual reconciliation of financial data by engineering a daily **ELT pipeline (Airbyte, dbt, BigQuery, Looker Studio)** that unified data across 3 disparate source systems for a **200+ project portfolio**, enabling real-time monitoring and reducing reporting errors for leadership.
- Built and deployed scalable **automated data pipelines** using **Python (Flask and FastAPI), Apache Airflow, Docker, dbt, and Databricks**, enabling near real-time analytics and improving **source file processing timeliness** and Power BI dashboard refresh performance.

- Implemented **automated testing, CI/CD pipelines** with GitHub Actions and Terraform, and version control standards, reducing pipeline failures and improving deployment reliability across data environments.
- Partnered with product managers, analysts, and program stakeholders using an **agile approach** to translate complex business requirements into scalable data architecture, including proofs-of-concept to demonstrate value before full implementation.

Hence AI | Data Engineer *May 2021 – Aug 2022*

- Built an **end-to-end data pipeline in Palantir Foundry** using Python, BeautifulSoup, and PySpark to ingest and transform web-scraped data, creating a **unified master view** of legal professional records for a curated marketplace.
- Designed Foundry ontologies modeling legal entities and applied **batch matching methodologies** to deduplicate and enrich records, enabling structured search and filtering across thousands of marketplace entries.
- Developed **PySpark transformations** to clean, deduplicate, and enrich source data using **fuzzy matching** and rule-based logic, improving dataset consistency, **data fitness of use**, and search accuracy across the platform.
- Implemented **ontology write-back workflows** to propagate user actions and maintain consistency between the data layer and underlying pipelines, enforcing **data quality rules** throughout the system.

Safaricom PLC | Software Engineer in Test *June 2019 — Aug 2020*

- Developed a **ReactJS dashboard** consuming REST APIs to visualize automation results and **analytical reporting** metrics, serving 20+ developers across 5 teams.
- Built a **pytest-based automated testing framework** for SOAP/REST APIs, implementing comprehensive unit, integration, and end-to-end tests to enforce data and service quality standards.
- Integrated testing frameworks into **Jenkins CI/CD pipelines**, automating test execution and ensuring continuous validation of microservices across environments.

PwC | Software Application Developer *Sep 2016 — Apr 2019*

- Owned data migration projects with **zero downtime**, demonstrating strong end-to-end delivery and data integrity management capabilities.
- Designed and implemented **ETL pipelines with SSIS** for batch data processing, cross-system integration, and **source file processing** across enterprise environments.
- Developed T-SQL stored procedures and applied **SQL-based data extraction and manipulation** across complex datasets to support analysis, quality monitoring, and reporting on MS SQL Server.
- Collaborated with cross-functional teams to gather requirements and deliver solutions aligned with business outcomes, including selection of tools and approaches for data projects.

Education

MSc Information Technology — Carnegie Mellon University
BSc Computer Science — University of Nairobi

Certifications

- AWS Certified Data Engineer Associate - April 2025
- AWS Certified Solutions Architect Associate - Aug 2021